

Protocol choices per stage and the largest effect measured here (SmolLM2-135M)

Calibration

- source
- window length
- count, seed

≤ 0.26 nats
seed SD 0.007



Quantization

- method, bits, group
- attention kernel
- threads

impl. ≤ 0.020
noise ≤ 0.004 nats



Evaluation

- eval set
- context, BOS
- dtype, tokenizer

≤ 0.021 nats



Reporting

- $D = \ln(\text{ppl} / \text{dense})$
- seeds, CIs, Holm
- hashes, config

comparable
effects



Latency: launcher (process vs inline), threads, model size — launcher up to $6.1\times$ (per token)