




(a) window-length effect



(b) SmolLM2-135M



(c) Qwen2.5-0.5B



Longer windows dilute the first-token (attention-sink) share of GPTQ's Hessian in sink-dominated down_proj layers. These layers hold ~39 % of SmolLM2's W3 GPTQ loss but ~1 % of Qwen's (REPORT.md A.6), matching the large vs absent effect in (a).